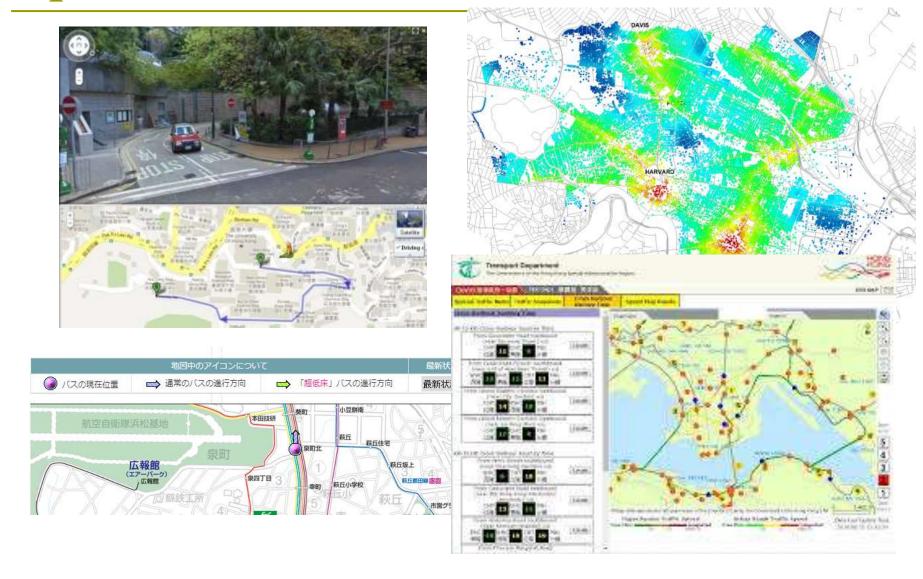
COMP7106B Big Data Management

Lecture 3 Spatial Networks

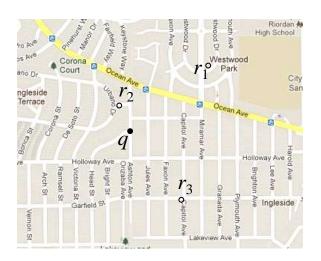
Prof. Reynold Cheng 24th February, 2024

Spatial Networks



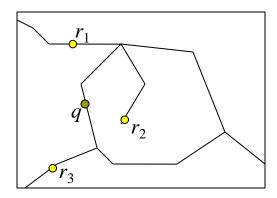
Spatial Networks

- Modeling and storing spatial networks
- Shortest path search
- Spatial queries over spatial networks
- Advanced indexing techniques for spatial networks



Network Distance

- In many real applications accessibility of objects is restricted by a spatial network
 - Examples
 - Driver looking for nearest gas station
 - Mobile user looking for nearest restaurant



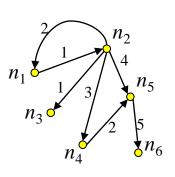
- Shortest path distance used instead of Euclidean distance
- SP(a,b) = path between a and b with the minimum accumulated length

Challenges

- Euclidean distance is no longer relevant
 - R-tree may not be useful, when search is based on shortest path distance
- Graph cannot be flattened to a onedimensional space
 - Special storage and indexing techniques for graphs are required
- Graph properties may vary
 - directed vs. undirected
 - length, time, etc. as edge weights

Modeling Spatial Networks

- Adjacency matrix only appropriate for dense graphs
- Spatial networks are sparse: use adjacency lists instead



graph

	n_1	n_2	n_3	n_4	n_5	n_6
n_1	0	1	8	8	8	8
n_2	2	0	1	3	4	8
n_3	8	8	0	8	8	8
n_4	8	8	8	0	2	8
n_5	8	8	8	8	0	5
n_6	8	8	8	8	8	0

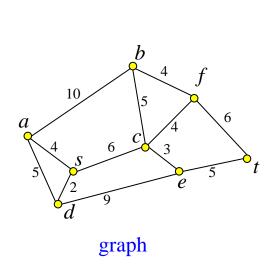
n_1	(n ₂ , 1)
n_2	$(n_1, 2), (n_3, 1), (n_4, 3), (n_5, 4)$
n_4	(n ₅ , 2)
n_5	(n ₆ , 5)

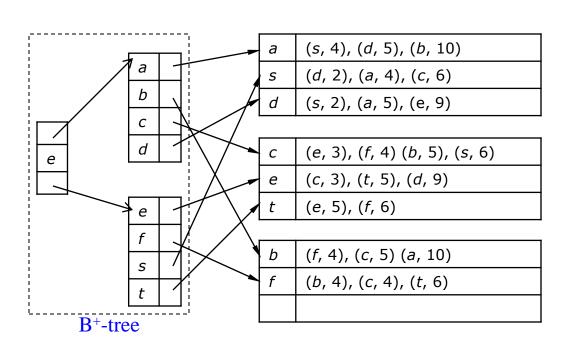
adjacency lists

adjacency matrix

Storing Large Spatial Networks

- Problem: adjacency lists representation may not fit in memory if graph is large
- Solution:
 - partition adjacency lists to disk blocks [based on proximity]
 - create B+-tree index on top of partitions [based on node-id]





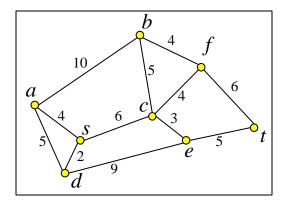
Shortest Path Computation

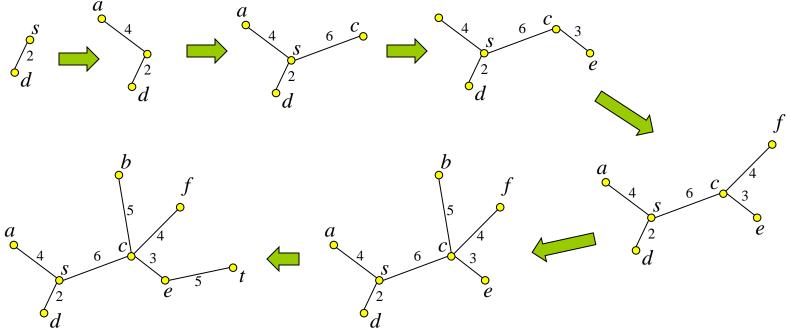
- Given a graph G(V,E), and two nodes s,t in V, find the shortest path from s to t
- A classic algorithmic problem
- Studied extensively since the 1950's
- Several methods
 - Dijkstra's algorithm
 - A*-search
 - Bi-directional search





Idea: incrementally explore the graph around s, visiting nodes in distance order to s until t is found (like NN)





Dijkstra's Shortest Path Search

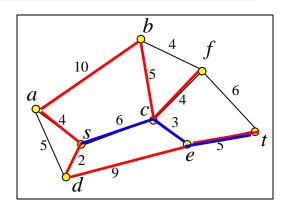
```
function Dijkstra\_SP(source node s, target node t)
1.
     for each graph node v
        SPD(s,v) := \infty; /* initialize shortest path distance */
2.
        path(s, v) := null; /* initialize shortest path */
3.
4.
        mark v as unvisited:
5.
     initialize a priority queue Q;
     SPD(s,s) := 0; add s to Q;
6.
7.
     while not empty(Q)
        v := top(Q); /* node v on Q with smallest SPD(s, v) */
8.
9.
        remove v from Q;
10.
        \max v as visited;
        if v = t then return path(s, t);
11.
        for each neighbor u of v
12.
           if u is not marked as visited
13.
14.
              if SPD(s, u) > SPD(s, v) + weight(v, u)
                 SPD(s, u) := SPD(s, v) + weight(v, u);
15.
                 path(s, u) := path(s, v) + (v, u);
16.
17.
                 add or update u on Q;
```

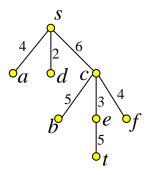
Dijkstra's Shortest Path Search

Current node	Queue
s (0)	d(2), a(4), c(6)
d (2)	a (4), c (6), e (11)
a (4)	c (6), e (11), b (14)
c (6)	e (9), f (10), b (11)
e (9)	f(10), b(11), t(14)
f(10)	b (11), t (14)
b (11)	t (14)

t is de-queued: shortest path has been found!

t (14)

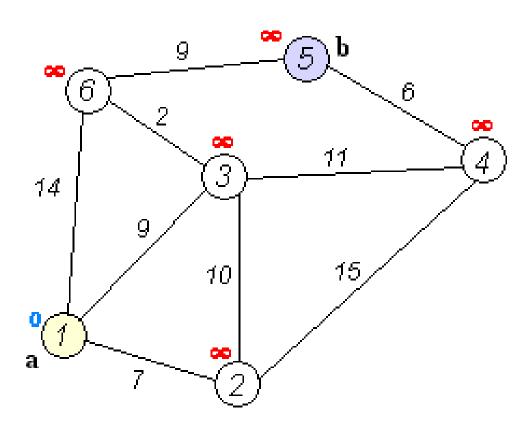




shortest path tree of node s

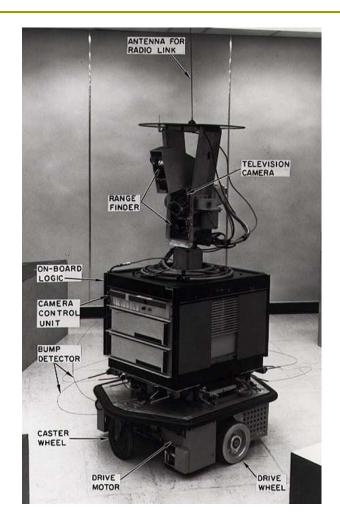
Illustrating Dijkstra's algorithm

- □ Find the shortest path between *a* and *b*.
- □ Worst-case performance $O(|E| + |V| \log |V|)$



A*-search

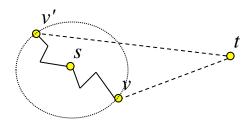
A* was invented by AI researchers working on Shakey the Robot's path planning in 1968.



https://en.wikipedia.org/wiki/A*_search_algorithm

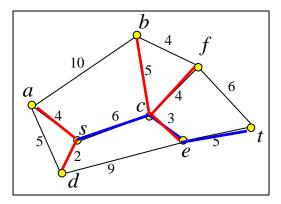
A*-search

- Dijkstra's search explores nodes around s without a specific search direction until t is found
- Idea: improve Dijkstra's algorithm by directing search towards t
- Due to triangular inequality, Euclidean distance is a lower bound of network distance
- Use Euclidean distance to lower bound network distance based on known information:
 - Nodes are visited in increasing SPD(s,v)+dist(v,t) order
 - □ SPD(s,v): shortest path distance from s to v (computed by Dijkstra)
 - \Box dist(v,t): Euclidean distance between v and t
 - Original Dijkstra visits nodes in increasing SPD(s,v) order



A*-search: Example

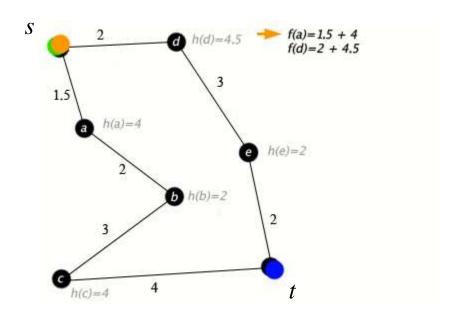
Current node	Queue $SPD(s,c)+dist(c,t)$
s (0)	d (2+14), a (4+15), c (6+7)
c (6)	a (4+15), d (2+14), b (11+9), f (10+6), e (9+5)
e (9)	a (4+15), d (2+14), b (11+9), f (10+6), t (14)
t (14)	



t is de-queued: shortest path has been found!

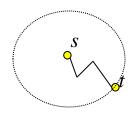
Illustrating A* search algorithm

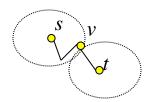
Find the shortest path between s and t.
f(p) = Dijkstra_dist(s, p) + Euclidean_dist(p, t)



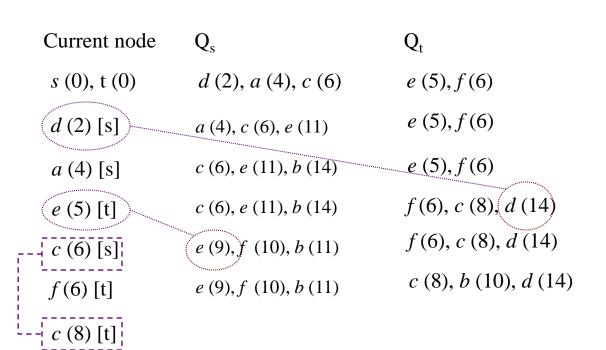
Bi-directional search

- Dijkstra's search explores nodes around s without a specific search direction until t is found
- Idea: search can be performed concurrently from s and from t (backwards)
- The shortest path tree of s and the (backward) shortest path tree of t are computed in concurrently
 - One queue Q_s for forward and one queue Q_t for backward search
 - Node visits are prioritized based on min(SPD(s,v), SPD(v,t))
 - If v already visited from s and v is in Q_t , then candidate shortest path: p(s,v)+p(v,t) [if v already visited from t and v in Q_s symmetric]
 - If v is visited by both s and t terminate search; report best candidate shortest path





Bi-directional search: Example



candidate shortest path: $s \rightarrow d \rightarrow e \rightarrow t (16)$

candidate shortest path: $s \rightarrow c \rightarrow e \rightarrow t (14)$

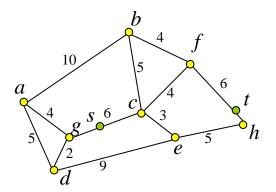
c is visited from both s and t! terminate and report shortest path

Discussions

- A* and bi-directional search can be combined to powerful search techniques
- A* can only be applied if lower distance bounds are available
- All versions of Dijkstra's search require nonnegative edge weights
 - Bellman-Ford is an algorithm for arbitrary negative edges

Source/Destination on Edges

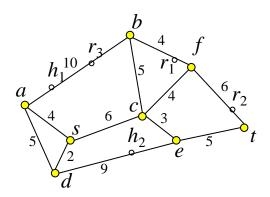
- We have assumed that points s and t are nodes of the network
- In practice s and t could be arbitrary points on edges
 - Mobile user locations
- Solve problem by introducing 2 more nodes

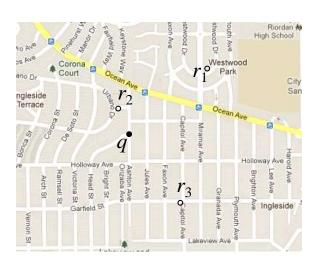


- en-heap g (2) and c (4) first
- t is reached from f or h

Spatial Queries over Spatial Networks

- Data:
 - A (static) spatial network (e.g., city map)
 - A (dynamic) set of spatial objects
- Spatial queries based on network distance:
 - Selections. Ex: find gas stations within 10km driving distance from here
 - Nearest neighbor search. Ex: find k nearest restaurants from present position
 - Joins. Ex: find pairs of restaurants and hotels at most 100m from each other



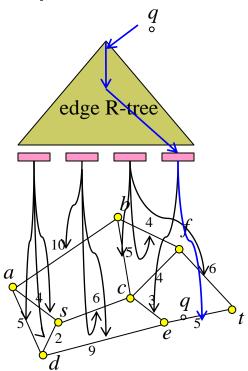


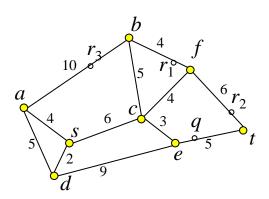
Methodology

- Store (and index) the spatial network
 - Graph component (indexes connectivity information)
 - Spatial component (indexes coordinates of nodes, edges, etc.)
- Store (and index) the sets of spatial objects
 - Ex., one spatial relation for restaurants, one spatial relation for hotels, one relation for mobile users, etc.
- Given a spatial location p, use spatial component of network to find the network edge containing p
- Given a network edge, use network component to traverse neighboring edges
- Given a neighboring edge, use spatial indexes to find objects on them

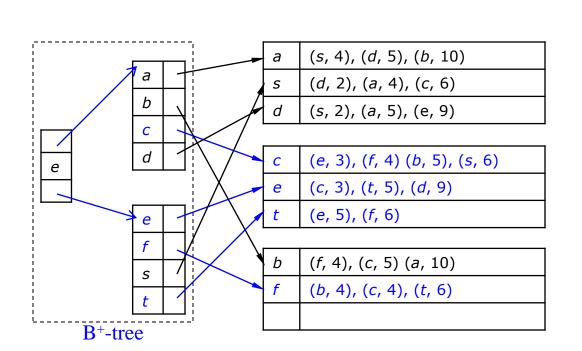
- Query: find all objects in spatial relation R, within network distance ε from location q
- Method:
 - Use spatial index of network (R-tree indexing network edges) to find edge n_1n_2 , which includes q
 - Use adjacency index of network (graph component) and apply Dijkstra's algorithm to progressively retrieve edges that are within network distance ε from location q
 - For all these edges apply a spatial selection on the Rtree that indexes R to find the results

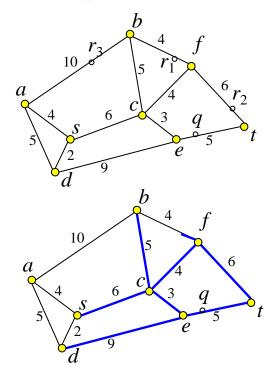
- Example: Find restaurants at most distance 10 from q
- Step 1: find network edge which contains q



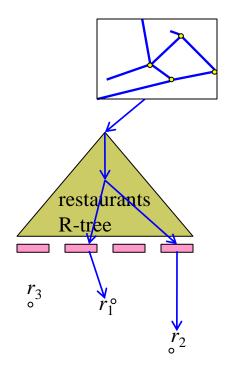


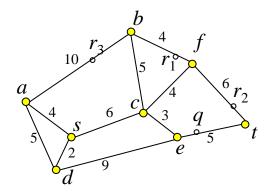
- Example: Find restaurants at most distance 10 from q
- Step 2: traverse network to find all edges (or parts of them within distance 10 from q)





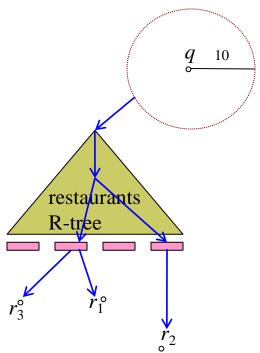
- Example: Find restaurants at most distance 10 from q
- Step 3: find restaurants that intersect the subnetwork computed at step 2

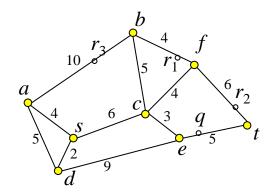




- Query: find all objects in spatial relation R, within network distance ε from location q
- Alternative method based on Euclidean bounds:
 - Assumption: Euclidean distance is a lower-bound of network distance:
 - □ dist(v,u) \le SPD(v,u), for any v,u
 - Use R-tree on R to find set S of objects such that for each o in S: dist(q,o) ≤ ε
 - For each o in S:
 - find where o is located in the network (use Network R-tree)
 - compute SPD(q,o) (e.g. use A*)
 - □ If SPD(q,o) \leq ε then output o

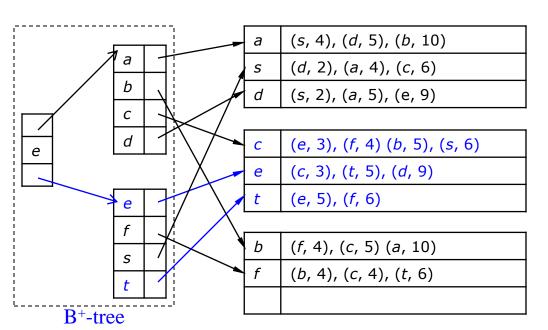
- Example: Find restaurants at most distance 10 from q
- Step 1: find restaurants for which the Euclidean distance to q is at most 10: $S = \{r_1, r_2, r_3\}$

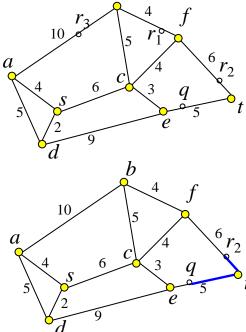




- Example: Find restaurants at most distance 10 from q
- Step 2: for each restaurant in S, compute SPD to q and verify if it is indeed a correct result

• Ex. for r_2 , first find where r_2 is located (edge f,t), then SP



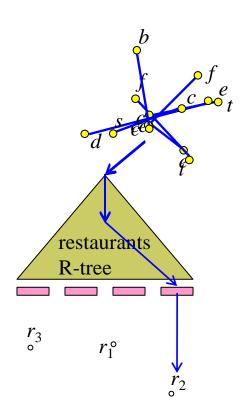


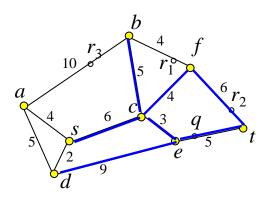
Evaluation of NN search (1)

- Query: find in spatial relation R the nearest object to a given location q
- Method:
 - Use spatial index of network (R-tree indexing network edges) to find edge n_1n_2 , which includes q
 - Use adjacency index of network (graph component) and apply Dijkstra's algorithm to progressively retrieve edges in order of their distance to q
 - For each edge apply a spatial selection on the R-tree that indexes R to find any objects
 - Keep track of nearest object found so far; use its shortest path distance to terminate network browsing

Evaluation of NN search (1)

Example: Find nearest restaurant to q





current
$$NN = r_2$$

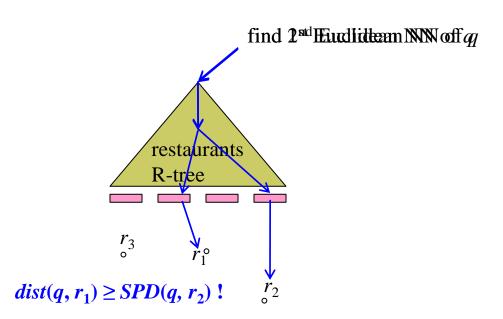
 $SPD(q, r_2)=5$

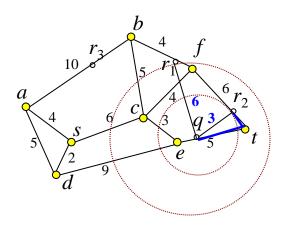
Evaluation of NN search (2)

- Query: find in spatial relation R the nearest object to a given location q
- Alternative method based on Euclidean bounds:
 - Assumption: Euclidean distance lower-bounds network distance:
 - □ dist(v,u) \le SPD(v,u), for any v,u
 - 1 Use R-tree on R to find Euclidean NN p_{F1} of q;
 - 2 CurrentNN= p_{E1} ; bound=SPD(q,p_{E1}); //e.g. use A*
 - 3 Find next Euclidean NN p_{Fi} of q
 - 4 If dist(q,Ei)≥bound, then report (CurrentNN,bound) as result;
 - 5 Compute SPD(q, p_{Ei}); if bound>SPD(q, p_{Ei}) then CurrentNN= p_{Ei} ; bound=SPD(q, p_{Ei});
 - 6 Goto step 3

Evaluation of NN search (2)

Example: Find nearest restaurant to q





current
$$NN = r_2$$

 $SPD(q, r_2) = 5$

Spatial Join Queries

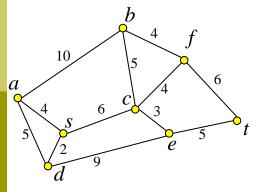
- Query: find pairs (r,s), such that r in relation R, s in relation S, and SPD(r,s)≤ε
- Methods:
 - For each r in R, do an ε-distance selection queries for objects in S (Index Nested Loops)
 - For each pair (r,s), such that Euclidean dist(r,s)≤ε compute SPD(r,s) and verify SPD(r,s)≤ε

Notes on Query Evaluation based on Network Distance

- For each query type, there are methods based on network browsing and methods based on Euclidean bounds
- Network browsing methods are fast if network edges are densely populated with points of interest
 - A limited network traversal can find the result fast
- Methods based on Euclidean bounds are good if the searched POIs are sparsely distributed in the network
 - Few verifications with exact SP searches are required
 - Directed SP search (e.g. using A*) avoids visiting empty parts of the network

Shortest Path Materialization and Indexing in Large Graphs

- Dijkstra's algorithm and related methods could be very expensive on very large graphs
- (Partial) materialization of shortest paths in static graphs can accelerate search



	а	b	С	S	
а	0, -	10, ab	10, asc	4, as	
b	10, ba	0, -	5, <i>bc</i>	11, bcs	:
С	10, csa	5, <i>cb</i>	0, -	6, <i>cs</i>	:
S	4, sa	11, scb	6, sc	0, -	

	а	b	с	S	
а	0, -	10, b	10, s	4, s	
b	10, a	0, -	5, c	11, c	•
С	10, s	5, b	0, -	6, <i>s</i>	•
s	4, a	11, c	6, <i>c</i>	0, -	

graph

brute-force materialization $O(n^3)$ space, O(1) time

distance matrix with successors $O(n^2)$ space, O(n) time

Summary

- Indexing and search of spatial networks is different than spatial indexing
 - Shortest path distance is used instead of Euclidean distance, to define range queries, nearest neighbor search, and spatial joins
- Spatial networks could be too large to fit in memory
 - Disk-based index for adjacency lists is used
- Several shortest path algorithms
- Spatial queries can be evaluated using Euclidean bounds