

Most materials in this course are based on materials from <http://aima.cs.berkeley.edu/> and <http://ai.berkeley.edu>



# 0. Introduction

Dirk Schnieders

# Intelligence

- We call ourselves Homo Sapiens
  - Latin: Man the wise
  - Intelligence is important to us
- For thousands of years, we have tried to understand how we humans think
  - How can our brain perceive, understand, predict, and manipulate a world far larger than itself ?
- Intelligence is most widely studied in humans, but has also been observed in animals and in plants

# Artificial Intelligence


- AI goes further than just understanding intelligence
  - It attempts to build intelligent entities
  - Computing to act effectively and safely in a wide variety of novel situations
- AI: science, or engineering?

# Why study AI?

AI's impact will be “more than anything in the history of mankind.”



# Why study AI?




Physics



AI

Historically,  
researchers have pursued several different versions of AI


- Let's build machines that ...
  - Think Humanly
  - Act Humanly
  - Think Rationally
  - Act Rationally



# A. Think Humanly



- Cognitive Modelling Approach: How do we think?
  - Introspection
  - Psychological experiments
  - Brain imaging
- Cognitive science constructs precise and testable theories of the human mind
  - E.g., express a theory as a computer program and compare input-output behaviors to a human
  - If there is a match, some of the programs mechanism could also be operating in humans





# A. Think Humanly

- The human brain is one of the great mysteries of science
  - How does our brain process information?
- The brain consists of nerve cells (aka neurons) and the collection of these simple cells leads to thought, action and consciousness
- The recent development of functional magnetic resonance imaging (fMRI) provides neuroscientists with details of brain activities




# A. Think Humanly

- Brains and digital computers have somewhat different properties
- A crude comparison of the raw computational resources

	Supercomputer	Personal Computer	Human Brain
Computational units	$10^6$ GPUs + CPUs	8 CPU cores	$10^6$ columns
	$10^{15}$ transistors	$10^{10}$ transistors	$10^{11}$ neurons
Storage units	$10^{16}$ bytes RAM	$10^{10}$ bytes RAM	$10^{11}$ neurons
	$10^{17}$ bytes disk	$10^{12}$ bytes disk	$10^{14}$ synapses
Cycle time	$10^{-9}$ sec	$10^{-9}$ sec	$10^{-3}$ sec
Operations/sec	$10^{18}$	$10^{10}$	$10^{17}$

- Would we be able to achieve the brain's level of intelligence with a computer of unlimited capacity?




## B. Act Humanly

- Turing Test Approach: Imitation Game was designed to provide a definition of intelligence
- A computer passes the test if a human interrogator, after posing some questions, cannot tell whether the response come from a human or a computer



Human  
interrogator



## B. Act Humanly

- The underlying principles of intelligence are more important than to duplicate an exemplar
- Consider another field: Artificial Flight
  - The Wright brothers succeeded because they stopped imitating birds and started using wind tunnels and learn about aerodynamics
  - It was not their goal to make “machines that fly so exactly like pigeons that they can fool even other pigeons”



## B. Act Humanly

### ➤ Reverse Turing Test

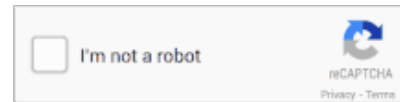
- Turing test in which the objective / roles have been reversed
- Interrogator is a computer. Interrogatee is human or computer.

### ➤ CAPTCHA

- Completely Automated Public Turing test to tell Computers and Humans Apart

following

finding



## C. Think Rationally

- The “laws of thought” approach
- What are the laws that guide and underlie our thinking?
- Greek schools developed various forms of logic
  - Notation and rules of derivation for thoughts
  - Example: Socrates is a man; all men are mortal; therefore, Socrates is mortal
- By 1965, programs existed that could (in principle) solve any solvable problem described in logic notation
- Problems with this approach
  - How to take informal knowledge and state it in formal terms?  
How about uncertainty?

# D. Act Rationally

- The rational agent approach
  - Act to achieve the best outcome
    - With uncertainty: the best expected outcome
- Advantages over the other approaches
  - More general
    - Correct inference is just one of several possible mechanisms for achieving rationality
  - Rationality is well defined
    - Human behaviour is well adapted only for one specific environment
- Most researchers in AI focus on the general principles of rational agents and how to build them

# Rational Agent

## ➤ Agent

- agent comes from the Latin *agere*, to do
- Something that perceives and acts
- E.g., robot or softbot

## ➤ Rational Agent

- Acts to achieve the best outcome or, when there is uncertainty, the best expected outcome
- For any given class of environments and tasks, we seek the agent (or class of agents) with the best performance



# Rational Agent

- AI focuses on the study and construction of rational agents
  - Agents that do the right thing
  - What counts as the right thing is defined by the objective that we provide to the agent
- Like other areas of research
  - Control Theory
  - Operations Research
  - Statistics
  - Economics
- This is called the standard model

# Perfect Rationality

- Always taking the exactly optimal action is not feasible in complex environments

# Issues with the Standard Model

- The standard model assumes that we will supply a fully specified objective function
  - Difficult in practice
- Value alignment problem
  - Example
    - Domain: Self-driving car
    - Objective: Reach destination safely
    - Problem
      - Strict goal of safety requires staying in the garage
      - There is a tradeoff between making progress towards the destination and incurring a risk of injury
      - How should this tradeoff be made?

# Value Alignment Problem

## ➤ Example: Chess



# Machine Learning

- An agent is learning if it improves its performance on future tasks
- Why would we want an agent to learn?
  - If the design of an agent can be improved, why not design the agent with that improvement to begin with?

# The Thinking Machine - MIT 1961



<http://www.youtube.com/watch?v=aygSMgK3BEM>

# History of AI - Turing Award Winners

- Marvin Minsky (1969)
- John McCarthy (1971)
- Edward Feigenbaum and Raj Reddy (1994)
- Judea Pearl (2011)
- Yoshua Bengio, Geoffrey Hinton, and Yann LeCun (2018)

# History of AI - Milestones

- Inception (1943 - 1956)
- Early Enthusiasm (1952 - 1969)
- A dose of reality (1966 - 1973)
- Expert systems (1969 - 1986)
- Return of NN (1986 - present)
- Probabilistic reasoning (1987 - present)
- Big data (2001 - present)
- Deep Learning (2011 - present)



# The State of the Art

## ➤ Publications

- AI papers increased 20 fold between 2010 to 2019 to 20,000 a year

## ➤ Conferences

- Attendance of [NeurIPS](#) increased 800% since 2012 to 13,500

## ➤ Industry

- AI start-ups in the US increased 20 fold from 2010 to 2019

## ➤ Internationalization (in 2019)

- China publishes more AI papers per year than US and about as many as Europe
- In citation weighted impact, US is ahead by 50% vs. China

# The State of the Art

## ➤ Vision

- Error rates for object detection improved from 28% to less than 2%

## ➤ Speed

- Training time for image recognition dropped by a factor of 100 in last 2 years
- Amount of computing power used in top AI applications is doubling every few month

## ➤ Humans vs. AI (in 2019)

- AI is better in chess, go, poker, pac-man, jeopardy!, object detection, speech recognition in limited domain, chinese-to-english in restricted domain, Quake III, Dota 2, StarCraft II, many Atari games, Skin cancer detection, prostate cancer detection, protein folding, ...

# Benefits of AI

- First solve AI, then use AI to solve everything else.



Demis Hassabis, Google DeepMind

# Risks of AI

- Lethal autonomous weapons
- Surveillance
- Biased decision making
- Impact on employment
- Safety-critical applications
- Cybersecurity

# Risks of AI - Superhuman AI

- Most experts agree that we will eventually be able to create a superhuman AI
  - An intelligence that far surpasses human ability

# Risks of AI - The Gorilla Problem

- About seven million years ago, a now-extinct primate evolved
  - one branch led to gorillas
  - another to humans
- Today the gorillas are probably not too happy about the human branch
  - They have no control over their future



# Risks of AI - The Gorilla Problem

- If the gorilla problem is the result of developing AI then we should stop working on it
- If superhuman AI (aka [AGI](#)) were a black box from outer space, we should be careful in opening the box
  - But it is not, we design the AI systems
  - If AI does end up taking control, it would be a design failure
- We need to understand the source of potential failure
  - Philosophical foundations of AI
  - Maybe the most important area of AI research

# AI Experts on the AI Apocalypse

- Worried AI experts signed on [open letter in March 2023](#) asking all AI labs to immediately pause “giant AI experiments”
- Where do AI experts stand regarding the probability of AGI and the probability of disaster by AGI?
  - Find out here
    - [https://dirk.hk/ai/IEEESpectrumAug2023\\_1.jpg](https://dirk.hk/ai/IEEESpectrumAug2023_1.jpg)
    - [https://dirk.hk/ai/IEEESpectrumAug2023\\_2.jpg](https://dirk.hk/ai/IEEESpectrumAug2023_2.jpg)



# Reading

- AIMA: Chapter 1
  - 66 pages available here: [https://dirk.hk/ai/AIMA\\_Chapter\\_1.pdf](https://dirk.hk/ai/AIMA_Chapter_1.pdf)
    - Password: Schnieders
- AIMA: Chapter 2
  - 49 pages available here: [https://dirk.hk/ai/AIMA\\_Chapter\\_2.pdf](https://dirk.hk/ai/AIMA_Chapter_2.pdf)
    - Password: Dirk
- The Thinking Machine – MIT 1961
  - 53 min video available here: <https://dirk.hk/ai/TheThinkingMachine.html>